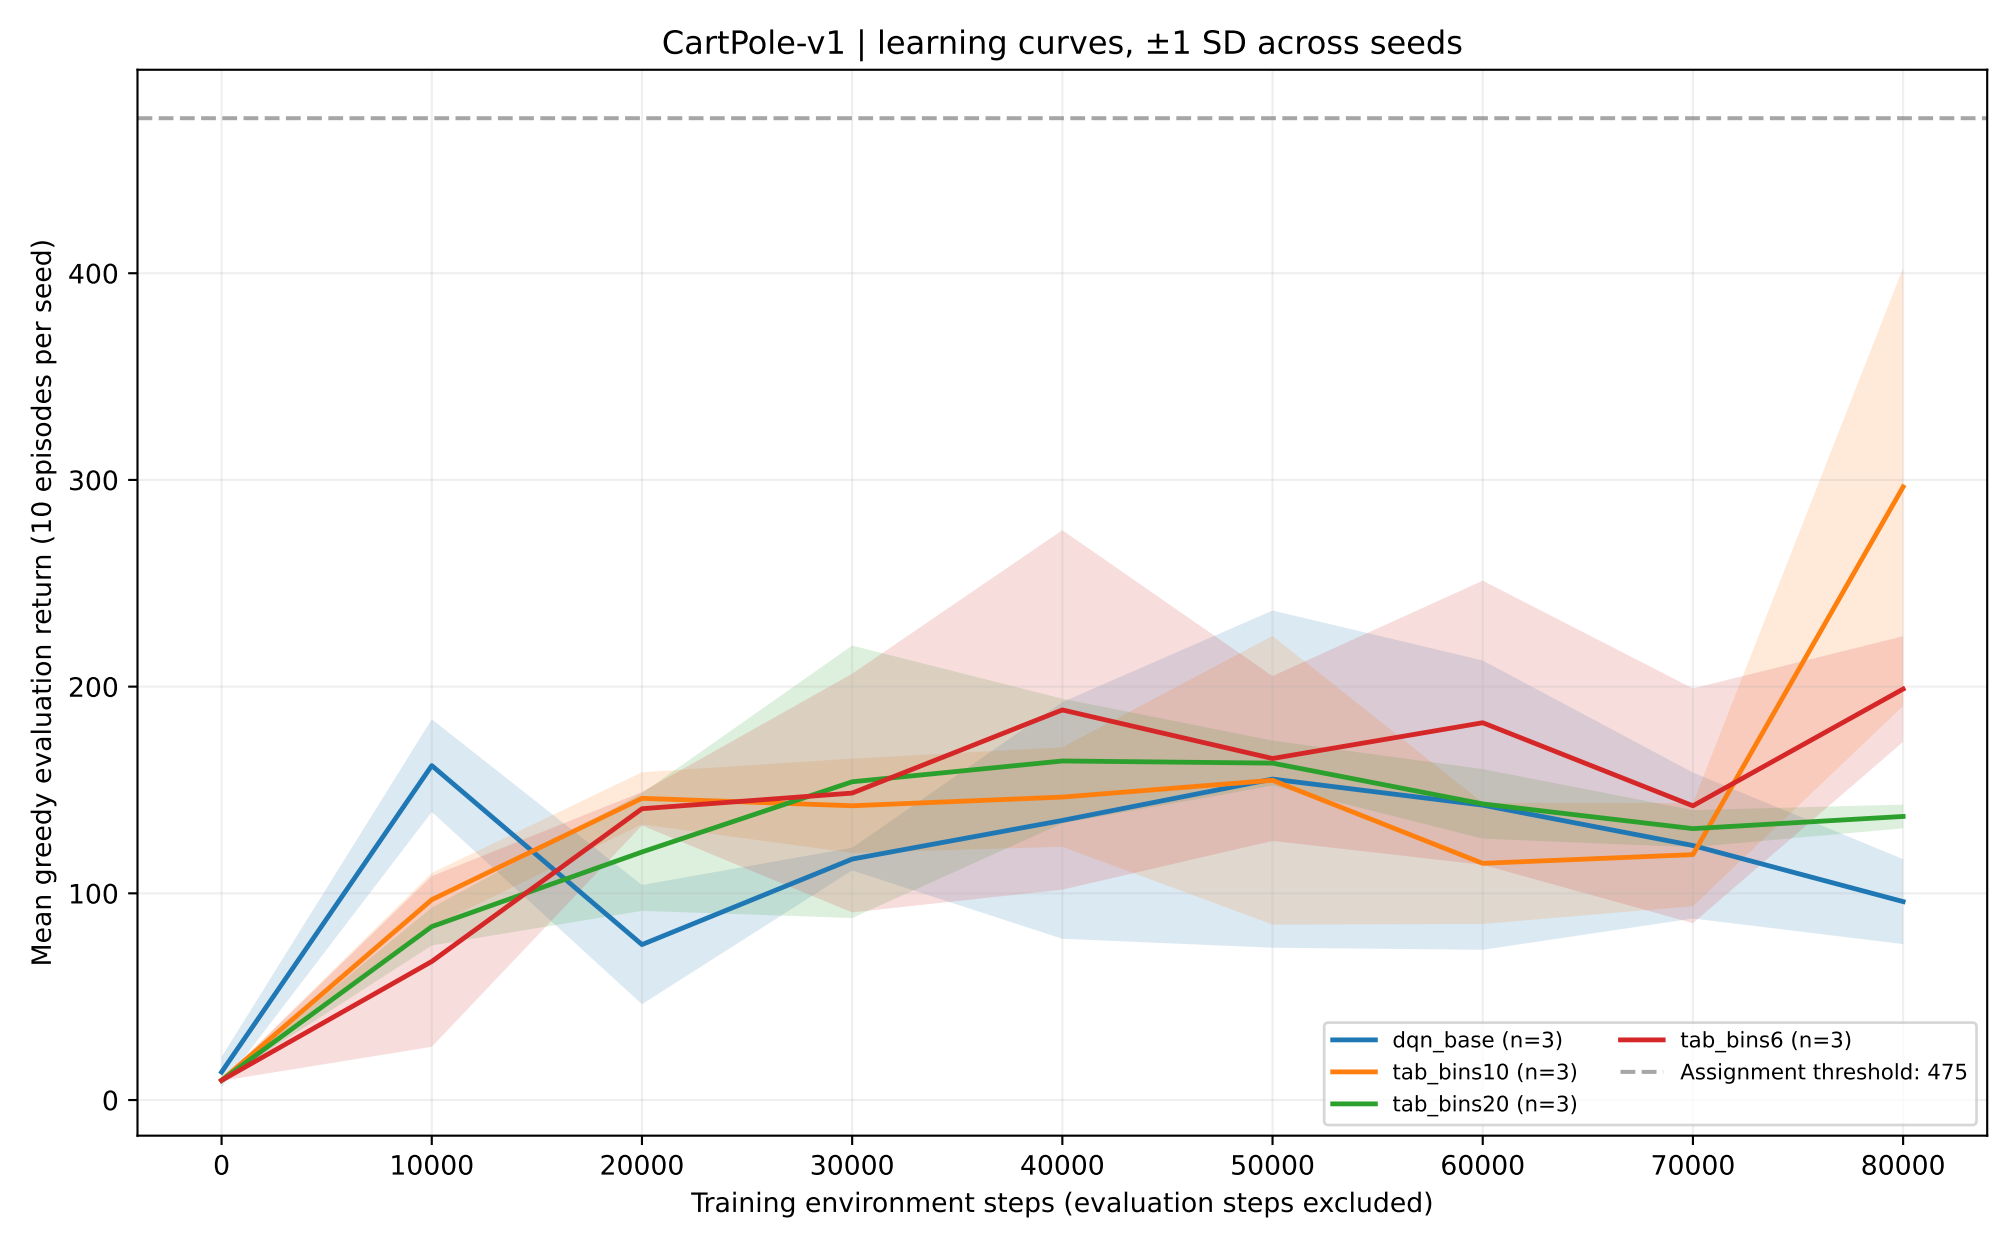
CartPole-v1 | learning curves, ±1 SD across seeds
Mean greedy evaluation return (10 episodes per seed)
0
100
200
300
400
0
10000
20000
30000
40000
50000
60000
70000
80000
Training environment steps (evaluation steps excluded)
dqn_base (n=3)
tab_bins10 (n=3)
tab_bins20 (n=3)
tab_bins6 (n=3)
Assignment threshold: 475