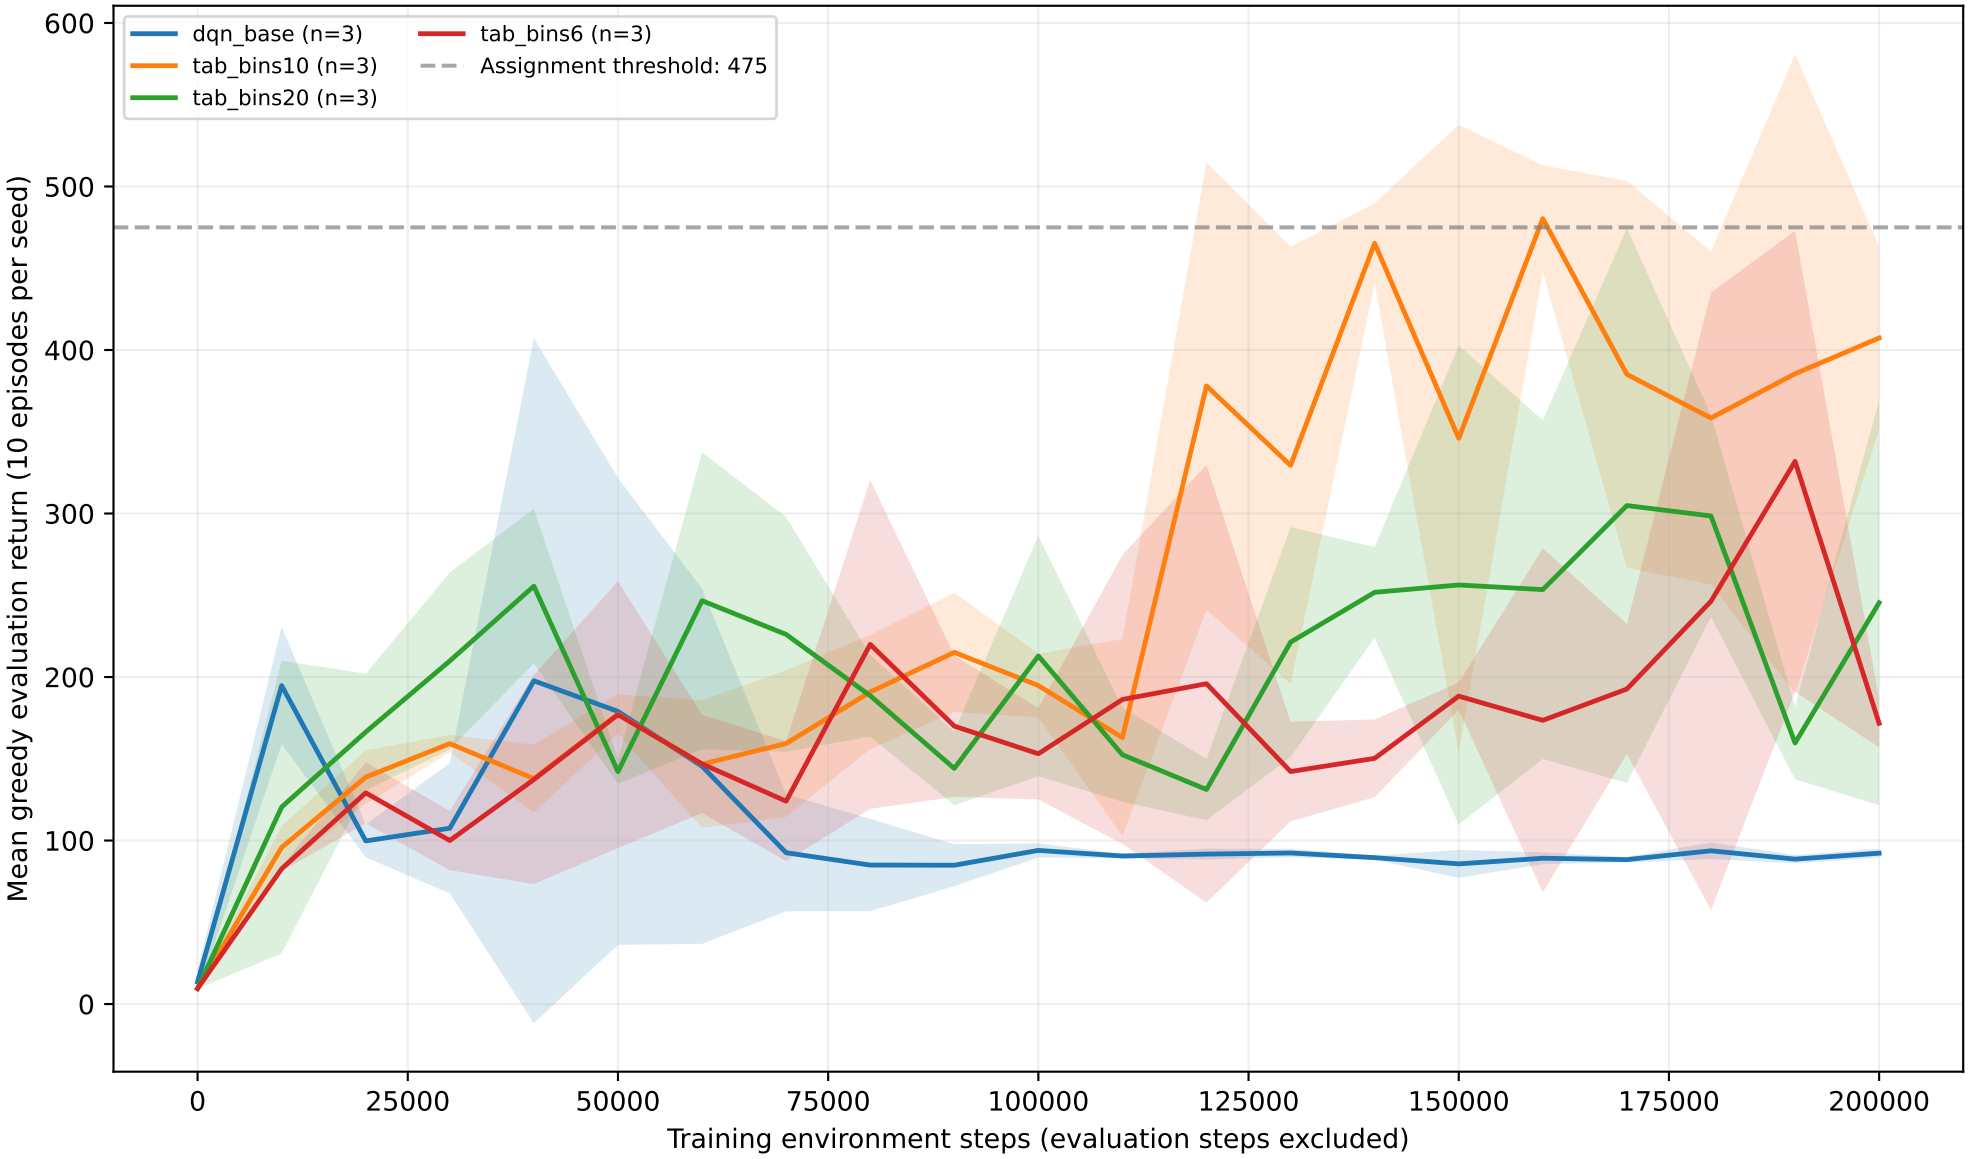
CartPole-v1 | learning curves, ±1 SD across seeds
dqn_base (n=3)
tab_bins10 (n=3)
tab_bins20 (n=3)
tab_bins6 (n=3)
Assignment threshold: 475
Mean greedy evaluation return (10 episodes per seed)
Training environment steps (evaluation steps excluded)
0
100
200
300
400
500
600
0
25000
50000
75000
100000
125000
150000
175000
200000